



香港中文大學 (深圳)
The Chinese University of Hong Kong

CSC3100 Data Structures


Lecture 20: Graph minimum spanning tree

Yixiang Fang
School of Data Science (SDS)
The Chinese University of Hong Kong, Shenzhen



Minimum spanning trees

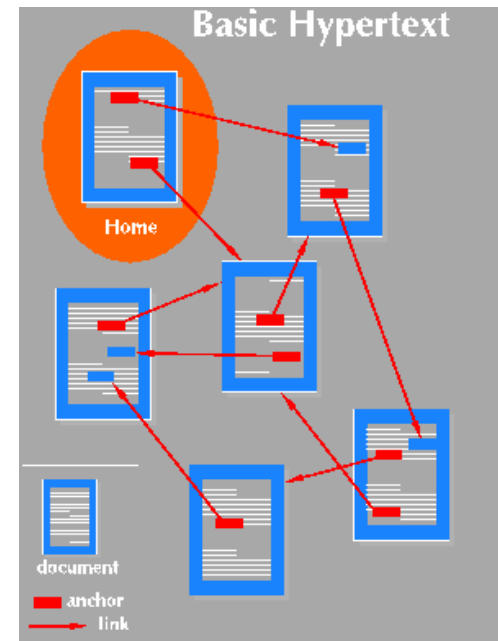

- ▶ **Spanning tree**
 - A tree (i.e., connected, acyclic graph) which contains all the vertices of the graph
- ▶ **Minimum spanning tree (MST)**
 - Spanning tree with the **minimum sum of weights**
 - If a graph is not connected, then there is an MST for each connected component of the graph





Applications of MST

- ▶ Find the least expensive way to connect a set of houses, cities, terminals, computers, etc.






Example

Problem

- ▶ A town has a set of houses and a set of roads
- ▶ A road connects 2 and only 2 houses
- ▶ A road connecting houses u and v has a cost $w(u, v)$



Goal: Build enough (and no more) roads such that:

1. Everyone stays connected
i.e., can reach every house from all other houses
2. Total cost is minimum




Minimum spanning trees (MSTs)

- ▶ A connected, undirected graph
 - Vertices = houses, Edges = roads
- ▶ A weighted $w(u, v)$ on each edge $(u, v) \in E$

Find $T \subseteq E$ such that:

1. T connects all vertices
2. $w(T) = \sum_{(u,v) \in T} w(u, v)$ is minimized




Properties of MST:

- (1) MST is **not** unique; (2) MST has no cycles;



Growing an MST: generic approach

- Grow a set A of edges (initially empty)
- Incrementally add edges to A such that they would belong to an MST




Idea: add only “safe” edges

- An edge (u, v) is **safe** for A , if and only if $A \cup \{(u, v)\}$ is also a subset of **some** MST



Generic MST algorithm

1. $A \leftarrow \emptyset$
2. **while** A is not a spanning tree
3. **do** find an edge (u, v) that is **safe** for A
4. $A \leftarrow A \cup \{(u, v)\}$
5. **return** A




- How do we find safe edges?



Finding safe edges

- ▶ Let's look at edge (h, g)
 - Is it safe for A initially?




- ▶ Yes. Why?
 - Let $S \subset V$ be any set of vertices that includes h but not g (so that g is in $V - S$)
 - In any MST, there has to be one edge (at least) that connects S with $V - S$
 - Why not choose the edge with **minimum weight** (h, g) ?



Definitions

- ▶ A **cut** $(S, V - S)$ is a partition of vertices into disjoint sets S and $V - S$
- ▶ An edge **crosses** the cut $(S, V - S)$ if one endpoint is in S and the other in $V - S$
- ▶ A cut **respects** a set A of edges \Leftrightarrow no edge in A crosses the cut
- ▶ An edge is a **light edge** crossing a cut \Leftrightarrow its weight is minimum over all edges crossing the cut
 - Note that for a given cut, there can be > 1 light edges crossing it






Theorem

- ▶ Let A be a subset of some MST, $(S, V - S)$ be a **cut** that respects A , and (u, v) be a **light edge** crossing $(S, V - S)$. Then (u, v) is safe for A .

Proof:


- ▶ Let T be an MST that includes A
 - edges in A are shaded
- ▶ Case1: If T includes (u, v) , then it would be safe for A
- ▶ Case2: Suppose T does not include the edge (u, v)
 - **Idea**: construct another MST T' that includes $A + \{(u, v)\}$





Theorem: proof

- ▶ T contains a unique path p between u and v
- ▶ Path p must cross the cut $(S, V - S)$ at least once: let (x, y) be that edge
- ▶ Let's remove $(x, y) \Rightarrow$ breaks T into two components
- ▶ Adding (u, v) reconnects the components
$$T' = T - \{(x, y)\} + \{(u, v)\}$$






Theorem: proof

$$T' = T - \{(x, y)\} + \{(u, v)\}$$

Have to show that T' is an MST:


- ▶ (u, v) is a light edge
 $\Rightarrow w(u, v) \leq w(x, y)$
- ▶ $w(T') = w(T) - w(x, y) + w(u, v) \leq w(T)$
- ▶ Since T is a spanning tree
 $w(T) \leq w(T') \Rightarrow T'$ must be an MST as well





Prim's algorithm

- ▶ The edges in set A always form a single tree
- ▶ Starts from an arbitrary "root": $V_A = \{a\}$
- ▶ At each step:
 - Find a light edge crossing $(V_A, V - V_A)$
 - Add this edge to A
 - Repeat until the tree spans all vertices







How to find light edges quickly?

Use a priority queue Q :

- Contains vertices not yet included in the tree, i.e., $(V - V_A)$
 - $V_A = \{a\}$, $Q = \{b, c, d, e, f, g, h, i\}$



- We associate a key with each vertex v :
 $\text{key}[v] = \text{minimum weight of any edge } (u, v) \text{ connecting } v \text{ to } V_A$






How to find light edges quickly?

- ▶ After adding a new node to V_A we update the weights of all the nodes adjacent to it



e.g., after adding a to the tree, $k[b]=4$ and $k[h]=8$

- ▶ Key of v is ∞ , if v is not adjacent to any vertices in V_A





Example




0 ∞ ∞ ∞ ∞ ∞ ∞ ∞ ∞ ∞
 $Q = \{a, b, c, d, e, f, g, h, i\}$
 $V_A = \emptyset$
Extract-MIN(Q) $\Rightarrow a$

key [b] = 4 π [b] = a
key [h] = 8 π [h] = a

4 ∞ ∞ ∞ ∞ ∞ ∞ **8** ∞
 $Q = \{b, c, d, e, f, g, h, i\}$ $V_A = \{a\}$
Extract-MIN(Q) $\Rightarrow b$



Example




key [c] = 8 π [c] = b
key [h] = 8 π [h] = a - unchanged

8 ∞ ∞ ∞ ∞ ∞ **8** ∞

$Q = \{c, d, e, f, g, h, i\}$ $V_A = \{a, b\}$

Extract-MIN(Q) \Rightarrow c



key [d] = 7 π [d] = c

key [f] = 4 π [f] = c

key [i] = 2 π [i] = c


7 ∞ **4** ∞ **8** **2**

$Q = \{d, e, f, g, h, i\}$ $V_A = \{a, b, c\}$

Extract-MIN(Q) \Rightarrow i



Example




key [h] = 7 π [h] = i

key [g] = 6 π [g] = i

7 ∞ 4 6 8

$Q = \{d, e, f, g, h\}$ $V_A = \{a, b, c, i\}$

Extract-MIN(Q) \Rightarrow f



key [g] = 2 π [g] = f

key [d] = 7 π [d] = c unchanged

key [e] = 10 π [e] = f


7 10 2 8

$Q = \{d, e, g, h\}$ $V_A = \{a, b, c, i, f\}$

Extract-MIN(Q) \Rightarrow g



Example




key [h] = 1 $\pi[h] = g$

7 10 1

$Q = \{d, e, h\}$ $V_A = \{a, b, c, i, f, g\}$

Extract-MIN(Q) \Rightarrow h




7 10

$Q = \{d, e\}$ $V_A = \{a, b, c, i, f, g, h\}$

Extract-MIN(Q) \Rightarrow d



Example



$\text{key}[e] = 9$ $\pi[e] = d$
9

$Q = \{e\}$ $V_A = \{a, b, c, i, f, g, h, d\}$

$\text{Extract-MIN}(Q) \Rightarrow e$

$Q = \emptyset$ $V_A = \{a, b, c, i, f, g, h, d, e\}$



PRIM(V, E, w, r)

```

1.  Q ← ∅
2.  for each u ∈ V
3.      do key[u] ← ∞
4.      π[u] ← NIL
5.      INSERT(Q, u)
6.  DECREASE-KEY(Q, r, 0)
7.  while Q ≠ ∅
8.      do u ← EXTRACT-MIN(Q)
9.          for each v ∈ Adj[u]
10.             do if v ∈ Q and w(u, v) < key[v]
11.                 then π[v] ← u
12.                     DECREASE-KEY(Q, v, w(u, v))

```

Total time: $O(V \lg V + E \lg V) = O(E \lg V)$
 $O(V)$ if Q is implemented as a min-heap

▶ key[r] ← 0 ← $O(\lg V)$
 Executed |V| times
 Min-heap operations: $O(V \lg V)$

Executed $O(E)$ times
 Constant
 Takes $O(\lg V)$
 $O(E \lg V)$



PRIM(V, E, w, r)

```

1.  Q ← ∅
2.  for each u ∈ V
3.      do key[u] ← ∞
4.          π[u] ← NIL
5.          INSERT(Q, u)
6.  DECREASE-KEY(Q, r, 0)
7.  while Q ≠ ∅
8.      do u ← EXTRACT-MIN(Q)
9.          for (j=0; j<|V|; j++)
10.             if (A[u][j]=1)
11.                 if v ∈ Q and w(u, v) < key[v]
12.                     then π[v] ← u
13.                     DECREASE-KEY(Q, v, w(u, v))

```

Total time: $O(V \lg V + E \lg V + V^2) = O(E \lg V + V^2)$

$O(V)$ if Q is implemented as a min-heap

▶ key[r] ← 0 ← $O(\lg V)$

Executed |V| times

Min-heap operations:

Takes $O(\lg V)$

Executed $O(V^2)$ times total

Constant


Takes $O(\lg V)$

$O(E \lg V)$




A different instance of the generic approach

(instance 1)



(instance 2)




- ▶ A is a forest containing connected components
 - Initially, each component is a single vertex
- ▶ Any safe edge merges two of these components into one
 - Each component is a tree



Kruskal's Algorithm


- ▶ How is it different from Prim's algorithm?
 - Prim's algorithm grows one tree all the time
 - Kruskal's algorithm grows multiple trees (i.e., a forest) at the same time
 - Trees are merged together using **safe** edges






Kruskal's Algorithm

- ▶ Start with each vertex being its own component
- ▶ Repeatedly merge two components into one by choosing the **light** edge that connects them
- ▶ Which components to consider at each iteration?
 - Scan the set of edges in monotonically increasing order by weight





Example



1: (h, g) 8: (a, h), (b, c)

2: (c, i), (g, f) 9: (d, e)

4: (a, b), (c, f) 10: (e, f)

6: (i, g) 11: (b, h)

7: (c, d), (i, h) 14: (d, f)

{a}, {b}, {c}, {d}, {e}, {f}, {g}, {h}, {i}

1. Add (h, g) {g, h}, {a}, {b}, {c}, {d}, {e}, {f}, {i}
2. Add (c, i) {g, h}, {c, i}, {a}, {b}, {d}, {e}, {f}
3. Add (g, f) {g, h, f}, {c, i}, {a}, {b}, {d}, {e}
4. Add (a, b) {g, h, f}, {c, i}, {a, b}, {d}, {e}
5. Add (c, f) {g, h, f, c, i}, {a, b}, {d}, {e}
6. Ignore (i, g) {g, h, f, c, i}, {a, b}, {d}, {e}
7. Add (c, d) {g, h, f, c, i, d}, {a, b}, {e}
8. Ignore (i, h) {g, h, f, c, i, d}, {a, b}, {e}
9. Add (a, h) {g, h, f, c, i, d, a, b}, {e}
10. Ignore (b, c) {g, h, f, c, i, d, a, b}, {e}
11. Add (d, e) {g, h, f, c, i, d, a, b, e}
12. Ignore (e, f) {g, h, f, c, i, d, a, b, e}
13. Ignore (b, h) {g, h, f, c, i, d, a, b, e}
14. Ignore (d, f) {g, h, f, c, i, d, a, b, e}



Algorithm 1: a straightforward method

Assume vertices are $1, 2, \dots, n$, and $E \geq V$

1. Sort all the edges $\longleftarrow O(E \log E)$
 2. **for** each $v \in V$
 3. $\text{map.add}(v, \{v\})$ $\left. \vphantom{\text{map.add}(v, \{v\})} \right\} O(V)$
 4. **for** each edge $(u, v) \in E$
 5. $\text{setU} = \text{map.get}(u), \text{setV} = \text{map.get}(v)$ $\left. \vphantom{\text{setU} = \text{map.get}(u)} \right\} O(1)$
 6. $\text{isConnected} = \text{false}$
 7. **for** each vertex $w \in \text{setU}$
 8. **if** $w == v$
 9. $\text{isConnected} = \text{true}$
 10. **break**
 11. **if** $\text{isConnected} == \text{false}$
 12. $\text{setU} = \text{setU} \cup \text{setV}$
 13. $\text{map.add}(u, \text{setU}), \text{map.add}(v, \text{setU})$ $\left. \vphantom{\text{map.add}(u, \text{setU})} \right\} O(\text{setV.size})$
 14. $R = R \cup \{(u, v)\}$
 15. Output R
- $\left. \vphantom{\text{map.add}(u, \text{setU})} \right\} O(VE)$

Can we do better?



Algorithm 2: using labels

- ▶ Using labels
 - A label means a connected component
 - Assign a unique label to each vertex initially
 - When merging two connected components, we always change the labels of vertices in the small component to the label of the large component
 - The cost of changing labels is smaller



Algorithm 2: using labels

Assume vertices are 1, 2, ..., n, and $E \supseteq V$

```
1.  Sort all the edges ←  $O(E \log E)$ 
2.  for each  $v \in V$ 
3.       $\text{label}[v] = v$ 
4.       $\text{setArray}[v] = \{v\}$ 
5.  for each edge  $(u, v) \in E$ 
6.       $uL = \text{label}[u], vL = \text{label}[v]$ 
7.      if  $uL == vL$  continue
8.       $R.\text{add}((u, v))$ 
9.      if  $\text{setArray}[uL].\text{size} \geq \text{setArray}[vL].\text{size}$ 
10.         for each vertex  $w \in \text{setArray}[vL]$ 
11.              $\text{label}[w] = uL$ 
12.              $\text{setArray}[uL].\text{add}(w)$ 
13.         else
14.             for each vertex  $w \in \text{setArray}[uL]$ 
15.                  $\text{label}[w] = vL$ 
16.                  $\text{setArray}[vL].\text{add}(w)$ 
17.  Output R
```

$O(V)$

$O(E)$

$O(\text{setArray}[vL].\text{size})$

$O(\text{setArray}[uL].\text{size})$


???



What's the total times of changing the labels for all the vertices?


▶ Case 1

- Height: $V-1$
- $O(V)$



▶ Case 2

- Height: $\lg V$
- $O(V \lg V)$



Consider a specific vertex v : if v 's label is changed, then the updated set $\text{setArray}[vL]$ will be at least twice larger than the original set $\text{setArray}[vL]$. Hence, the number of times for changing labels for v is at most $O(\lg V)$.



Algorithm 2: using labels

Assume vertices are 1, 2, ..., n, and $E \geq V$

```
1.  Sort all the edges ←  $O(E \log E)$ 
2.  for each  $v \in V$ 
3.      label[v] = v
4.      setArray[v] = {v} }  $O(V)$ 
5.  for each edge  $(u, v) \in E$ 
6.       $uL = \text{label}[u], vL = \text{label}[v]$ 
7.      if  $uL == vL$  continue }  $O(E)$ 
8.      R.add((u, v))
9.      if setArray[uL].size  $\geq$  setArray[vL].size
10.         for each vertex  $w \in \text{setArray}[vL]$ 
11.             label[w] = uL
12.             setArray[uL].add(w)
13.         else
14.             for each vertex  $w \in \text{setArray}[uL]$ 
15.                 label[w] = vL
16.                 setArray[vL].add(w)
17.  Output R
```

$O(E \log E)$

$O(V \log V)$



Recommended reading

- ▶ Reading materials
 - Textbook Chapter 23
- ▶ Next lecture
 - Shortest paths, Chapters 24&25